

DevOps Interview Preparation (Quick Glance)

AWS

Check Availability Zone Availability

```
aws ec2 describe-instance-type-offerings \
  --location-type availability-zone \
  --filters "Name=instance-
type,Values=$instance_type"
"Name=location,Values=$1" \
  --region $region \
  --query
'InstanceTypeOfferings[?InstanceType=='${ins
tance_type}'].Location' \
  --output text
```

Key Pairs

Check if Key Pair Exists:

```
if ! aws ec2 describe-key-pairs --key-names
${key_pair_name} --region ${region}
&>/dev/null; then
```

Create Key Pair:

```
aws ec2 create-key-pair --key-name
${key_pair_name} --query 'KeyMaterial' --
output text --region ${region} >
CentosComplexKeyPair.pem
```

Set Permissions:

```
chmod 400 CentosComplexKeyPair.pem
```

VPC

Describe VPCs:

```
aws ec2 describe-vpcs --filters
"Name=cidr,Values=${vpc_cidr}" --query
'Vpcs[0].VpcId' --output text --region
${region}
```

Create VPC:

```
aws ec2 create-vpc --cidr-block ${vpc_cidr} -
-query 'Vpc.VpcId' --output text --region
${region}
```

Internet Gateway

Check if Internet Gateway Exists:

```
igw_id=$(aws ec2 describe-internet-gateways -
-filters "Name=attachment.vpc-
id,Values=${vpc_id}" --query
'InternetGateways[0].InternetGatewayId' --
output text --region ${region})
if [ "$igw_id" == "None" ]; then
```

Create Internet Gateway:

```
igw_id=$(aws ec2 create-internet-gateway --
query 'InternetGateway.InternetGatewayId' --
output text --region ${region})
```

Attach Internet Gateway:

```
aws ec2 attach-internet-gateway --internet-
gateway-id ${igw_id} --vpc-id ${vpc_id} --
region ${region}
```

Subnets

Check if Public Subnet 1 Exists:

```
public_subnet_id_1=$(aws ec2 describe-subnets
--filters "Name=vpc-id,Values=${vpc_id}"
"Name=cidr-
block,Values=${public_subnet_cidr_1}" --query
'Subnets[0].SubnetId' --output text --region
${region})
if [ "$public_subnet_id_1" == "None" ]; then
```

Create Public Subnet 1:

```
public_subnet_id_1=$(aws ec2 create-subnet --
vpc-id ${vpc_id} --cidr-block
${public_subnet_cidr_1} --availability-zone
${available_zone_1} --query 'Subnet.SubnetId'
--output text --region ${region})
```

Route Tables

Check if Route Table for Public Subnet 1 Exists:

```
public_route_table_id_1=$(aws ec2 describe-
route-tables --filters "Name=vpc-
id,Values=${vpc_id}"
"Name=association.subnet-
id,Values=${public_subnet_id_1}" --query
```

```
'RouteTables[0].RouteTableId' --output text -
-region ${region})
if [ "$public_route_table_id_1" == "None" ];
then
```

Create Route Table for Public Subnet 1:

```
public_route_table_id_1=$(aws ec2 create-
route-table --vpc-id ${vpc_id} --query
'RouteTable.RouteTableId' --output text --
region ${region})
```

Associate Route Table with Public Subnet 1:

```
aws ec2 associate-route-table --route-table-
id ${public_route_table_id_1} --subnet-id
${public_subnet_id_1} --region ${region}
```

Create Route in Route Table for Public Subnet 1:

```
aws ec2 create-route --route-table-id
${public_route_table_id_1} --destination-
cidr-block 0.0.0.0/0 --gateway-id ${igw_id} -
-region ${region}
```

NAT Gateway

Allocate Elastic IP:

```
eip_allocation_id_1=$(aws ec2 allocate-
address --domain vpc --query 'AllocationId' -
-output text --region ${region})
```

Create NAT Gateway:

```
nat_gateway_id_1=$(aws ec2 create-nat-gateway
--subnet-id ${public_subnet_id_1} --
allocation-id ${eip_allocation_id_1} --query
'NatGateway.NatGatewayId' --output text --
region ${region})
```

Update Private Route Table 1:

```
aws ec2 create-route --route-table-id
${private_route_table_id_1} --destination-
cidr-block 0.0.0.0/0 --nat-gateway-id
${nat_gateway_id_1} --region ${region}
echo "Updated Private Route Table 1 to use
NAT Gateway 1"
```

Security Groups

Check if Bastion Security Group Exists:

```
bastion_security_group_id=$(aws ec2 describe-
security-groups --filters "Name=vpc-
id,Values=${vpc_id}" "Name=group-
name,Values=${bastion_security_group_name}" -
-query 'SecurityGroups[0].GroupId' --output
text --region ${region})
if [ "$bastion_security_group_id" == "None"
]; then
```

Create Bastion Security Group:

```
bastion_security_group_id=$(aws ec2 create-
security-group --group-name
${bastion_security_group_name} --description
"Bastion security group" --vpc-id ${vpc_id} -
-query 'GroupId' --output text --region
${region})
```

Add Inbound Rules to Bastion Security Group:

```
aws ec2 authorize-security-group-ingress --
group-id ${bastion_security_group_id} --
protocol tcp --port 22 --cidr 0.0.0.0/0 --
region ${region}
```

Check if Application Security Group Exists:

```
app_security_group_id=$(aws ec2 describe-
security-groups --filters "Name=vpc-
id,Values=${vpc_id}" "Name=group-
name,Values=${app_security_group_name}" --
query 'SecurityGroups[0].GroupId' --output
text --region ${region})
if [ "$app_security_group_id" == "None" ];
then
```

Create Application Security Group:

```
app_security_group_id=$(aws ec2 create-
security-group --group-name
${app_security_group_name} --description
"Application security group" --vpc-id
${vpc_id} --query 'GroupId' --output text --
region ${region})
```

Add Inbound Rules to Application Security Group:

```
aws ec2 authorize-security-group-ingress --group-id ${app_security_group_id} --protocol tcp --port 22 --source-group ${bastion_security_group_id} --region ${region}
aws ec2 authorize-security-group-ingress --group-id ${app_security_group_id} --protocol tcp --port 80 --cidr 0.0.0.0/0 --region ${region}
```

IAM Role

Trust Policy:

```
cat > trust-policy.json <<EOF
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "ec2.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    },
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "lambda.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
EOF
```

Create Role:

```
aws iam create-role --role-name ${role_name} --assume-role-policy-document file://trust-policy.json --region ${region}
```

Attach Policy:

```
aws iam attach-role-policy --role-name ${role_name} --policy-arn ${policy_arn} --region ${region}
```

Create Instance Profile:

```
aws iam create-instance-profile --instance-profile-name ${instance_profile_name} --region ${region}
```

Add Role to Instance Profile:

```
aws iam add-role-to-instance-profile --instance-profile-name ${instance_profile_name} --role-name ${role_name} --region ${region}
```

Launch EC2 Instance with Instance Profile:

```
aws ec2 run-instances --image-id ami-0abcdef1234567890 --count 1 --instance-type t2.micro --iam-instance-profile Name=MyInstanceProfile --region us-west-2
```

Placement Group

Create Placement Group:

```
aws ec2 create-placement-group --group-name ${placement_group_name} --strategy spread --region ${region}
```

- --group-name \${placement_group_name} : Specifies the name of the placement group.

strategy spread : Specifies the placement strategy (spread in this case).

- --region \${region} : Specifies the AWS region.

Cluster Placement Group

Use Case: High-performance computing (HPC) applications, big data workloads, and applications that require high network throughput.

```
aws ec2 create-placement-group --group-name my-cluster-group --strategy cluster --region us-west-2
```

Spread Placement Group

Use Case: Applications that require high availability and need to be isolated from failures, such as critical applications.

```
aws ec2 create-placement-group --group-name my-spread-group --strategy spread --region us-west-2
```

Partition Placement Group

Use Case: Large distributed and replicated workloads, such as Hadoop, Cassandra, and Kafka.

```
aws ec2 create-placement-group --group-name my-partition-group --strategy partition --partition-count 3 --region us-west-2
```

This command creates a partition placement group named my-partition-group with 3 partitions in the us-west-2 region.

Launch Instances in the Partition Placement Group:

```
aws ec2 run-instances --image-id ami-0abcdef1234567890 --count 3 --instance-type t2.micro --placement "GroupName=my-partition-group,PartitionNumber=0" --region us-west-2
aws ec2 run-instances --image-id ami-0abcdef1234567890 --count 3 --instance-type t2.micro --placement "GroupName=my-partition-group,PartitionNumber=1" --region us-west-2
aws ec2 run-instances --image-id ami-0abcdef1234567890 --count 3 --instance-type t2.micro --placement "GroupName=my-partition-group,PartitionNumber=2" --region us-west-2
```

S3 Bucket

Create S3 Bucket:

```
aws s3api create-bucket --bucket ${bucket_name} --region ${region} --create-bucket-configuration LocationConstraint=${region}
```

Create Sample File:

```
echo "This is a sample file for S3 bucket." > sample_file.txt
```

Upload Sample File:

```
aws s3 cp sample_file.txt s3://${bucket_name}/sample_file.txt --region ${region}
```

RDS

Create RDS Instance:

```
aws rds create-db-instance \
  --db-instance-identifier ${db_instance_identifier} \
  --db-instance-class ${db_instance_class} \
  --engine ${engine} \
  --master-username ${master_username} \
  --master-user-password ${master_user_password} \
  --allocated-storage 20 \
  --db-name ${db_name} \
  --vpc-security-group-ids ${app_security_group_id} \
  --db-subnet-group-name ${db_subnet_group_name} \
  --multi-az \
  --no-publicly-accessible \
  --region ${region}
```

Wait for Availability:

```
aws rds wait db-instance-available --db-instance-identifier ${db_instance_identifier} --region ${region}
```

Get RDS Endpoint:

```
db_endpoint=$(aws rds describe-db-instances -
-db-instance-identifier
${db_instance_identifier} --query
'DBInstances[0].Endpoint.Address' --output
text --region ${region})
echo "RDS instance endpoint: ${db_endpoint}"
```

Create DB Subnet Group:

```
aws rds create-db-subnet-group \
--db-subnet-group-name
${db_subnet_group_name} \
--db-subnet-group-description "My DB
Subnet Group" \
--subnet-ids ${private_subnet_id_1}
${private_subnet_id_2} \
--region ${region}
```

AWS CloudWatch

Create CloudWatch Alarm:

```
aws cloudwatch put-metric-alarm --alarm-name
${alarm_name} \
--metric-name CPUUtilization --namespace
AWS/EC2 \
--statistic Average --period 300 --
threshold 80 \
--comparison-operator
GreaterThanOrEqualToThreshold \
--dimensions
Name=InstanceId,Value=${instance_ids[0]} \
--evaluation-periods 2 --alarm-actions
${sns_topic_arn} \
--region ${region}
```

Launch Instances User Data Script:

```
cat > userDataCentOsComplex.sh <<EOF
#!/bin/bash
# Install httpd, unzip, and aws-cli
yum update -y
yum install -y httpd unzip aws-cli

# Start httpd service
systemctl start httpd
s
# Enable httpd service to start on boot
systemctl enable httpd

# Create a sample log file
echo "This is a sample log file." >
./sample_log.txt

# Upload the log file to S3 bucket
bucket_name=$(grep bucket_name
./resource_ids_centos.txt | cut -d'=' -f2)
aws s3 cp ./sample_log.txt
s3://${bucket_name}/sample_log.txt

# Download and unzip the website files
cd /var/www/html
wget
https://www.tooplate.com/download/2137_barist
a_cafe -O barista_cafe.zip
EOF
```

Launch Instances:

```
aws ec2 run-instances \
--image-id ami-0abcdef1234567890 \
--count 2 \
--instance-type t3.micro \
--key-name ${key_pair_name} \
--security-group-ids
${app_security_group_id} \
--subnet-id ${private_subnet_id_2} \
--user-data
file://userDataCentOsComplex.sh \
--tag-specifications
'ResourceType=instance,Tags=[{Key=Name,Value=
'${instance_name}_2'}]' \
--region ${region} \
--monitoring "Enabled=false" \
```

```
--iam-instance-profile
Name=${instance_profile_name} \
--block-device-mappings
'[{ "DeviceName": "/dev/sdh", "Ebs": { "VolumeSize
": 8, "DeleteOnTermination": true } }]' \
--placement
"AvailabilityZone=${available_zone_2},GroupNa
me=${placement_group_name}" \
--instance-initiated-shutdown-behavior
"terminate" \
--query 'Instances[*].InstanceId' --
output text
```

Wait for Running State:

```
aws ec2 wait instance-running --instance-ids
${instance_ids} --region ${region}
```

Wait for Status Checks to Pass:

```
aws ec2 wait instance-status-ok --instance-
ids ${instance_ids} --region ${region}
```

Load Balancers

Create Load Balancer:

```
load_balancer_arn=$(aws elbv2 create-load-
balancer \
--name my-load-balancer \
--subnets ${public_subnet_id_1}
${public_subnet_id_2} \
--security-groups
${app_security_group_id} \
--query
'LoadBalancers[0].LoadBalancerArn' --output
text --region ${region})
```

Create Target Group:

```
target_group_arn=$(aws elbv2 create-target-
group \
--name my-target-group \
--protocol HTTP \
--port 80 \
--vpc-id ${vpc_id} \
--query 'TargetGroups[0].TargetGroupArn'
--output text --region ${region})
```

AutoScaling Group

Create Launch Template:

```
launch_template_id=$(aws ec2 create-launch-
template \
--launch-template-name
${launch_template_name} \
--version-description "v1" \
--launch-template-data '{
"ImageId": "${image_id}",
"InstanceType": "t3.micro",
"KeyName": "${key_pair_name}",
"SecurityGroupIds":
["${app_security_group_id}"],
"IamInstanceProfile": {"Name":
"${instance_profile_name}"},
"UserData": "${(base64 -w 0
./userDataCentOsComplex.sh) }",
"BlockDeviceMappings": [{
"DeviceName": "/dev/sdh",
"Ebs": {
"VolumeSize": 8,
"DeleteOnTermination": true
}
}]}
}' --query
'LaunchTemplate.LaunchTemplateId' --output
text --region ${region})
```

Create Auto Scaling Group:

```
aws autoscaling create-auto-scaling-group \
--auto-scaling-group-name
${auto_scaling_group_name} \
--launch-template
"LaunchTemplateId=${launch_template_id},Versi
on=1" \
--min-size ${min_size} \
--max-size ${max_size} \
```

```
--desired-capacity ${desired_capacity} \
--vpc-zone-identifier "${subnet_ids}" \
--region ${region}
```

Scale Up Policy:

```
scale_up_policy_arn=$(aws autoscaling put-
scaling-policy \
  --auto-scaling-group-name
${auto_scaling_group_name} \
  --policy-name ScaleUpPolicy \
  --scaling-adjustment 1 \
  --adjustment-type ChangeInCapacity \
  --region ${region} \
  --query 'PolicyARN' --output text)
```

Scale Down Policy:

```
scale_down_policy_arn=$(aws autoscaling put-
scaling-policy \
  --auto-scaling-group-name
${auto_scaling_group_name} \
  --policy-name ScaleDownPolicy \
  --scaling-adjustment -1 \
  --adjustment-type ChangeInCapacity \
  --region ${region} \
  --query 'PolicyARN' --output text)
```

High CPU Utilization Alarm:

```
aws cloudwatch put-metric-alarm \
  --alarm-name HighCPUUtilization \
  --metric-name CPUUtilization \
  --namespace AWS/EC2 \
  --statistic Average \
  --period 300 \
  --threshold 80 \
  --comparison-operator
GreaterThanOrEqualToThreshold \
  --dimensions
Name=AutoScalingGroupName,Value=${auto_scalin
g_group_name} \
  --evaluation-periods 2 \
  --alarm-actions ${scale_up_policy_arn} \
  --region ${region}
```

Low CPU Utilization Alarm:

```
aws cloudwatch put-metric-alarm \
  --alarm-name LowCPUUtilization \
  --metric-name CPUUtilization \
  --namespace AWS/EC2 \
  --statistic Average \
  --period 300 \
  --threshold 20 \
  --comparison-operator
LessThanOrEqualToThreshold \
  --dimensions
Name=AutoScalingGroupName,Value=${auto_scalin
g_group_name} \
  --evaluation-periods 2 \
  --alarm-actions ${scale_down_policy_arn}
\
  --region ${region}
```

Kubernetes

Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  labels:
    app: my-app
spec:
  containers:
    - name: my-container
      image: nginx:1.14.2
      ports:
        - containerPort: 80
      resources:
        requests:
          cpu: "100m"
          memory: "128Mi"
        limits:
```

```
cpu: "500m"
memory: "256Mi"
readinessProbe:
  httpGet:
    path: /
    port: 80
    initialDelaySeconds: 5
    periodSeconds: 10
livenessProbe:
  httpGet:
    path: /healthz
    port: 80
    initialDelaySeconds: 15
    periodSeconds: 20
restartPolicy: Always
nodeSelector:
  disktype: ssd
tolerations:
- key: "key"
  operator: "Equal"
  value: "value"
  effect: "NoSchedule"
```

ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
  labels:
    app: my-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
  labels:
    app: my-app
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          imagePullPolicy: IfNotPresent
```

Service

```
apiVersion: v1
kind: Service
```

```

metadata:
  name: my-service
  annotations:
    service.beta.kubernetes.io/aws-load-
balancer-type: "nlb"
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      name: http
      type: LoadBalancer
      sessionAffinity: ClientIP
      externalTrafficPolicy: Local

```

ConfigMap

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config
data:
  config.property: "some-value"
  another.property: |
    line1
    line2
binaryData:
  binaryFile: <base64 encoded>

```

Secret

```

apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  username: dXNlcm5hbWU= # base64 encoded
  password: cGFzc3dvcmQ= # base64 encoded
stringData:
  config.yaml: |
    apiUrl: "https://myapi.com"
    token: "my-token"

```

PersistentVolume

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0001
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  storageClassName: standard
  nfs:
    server: nfs-server.example.com
    path: "/exports"

```

PersistentVolumeClaim

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: standard
  volumeMode: Filesystem
  volumeName: pv0001 # optional, binds to
a specific PV

```

Namespace

```

apiVersion: v1
kind: Namespace
metadata:

```

```

name: my-namespace
labels:
  environment: development

```

DaemonSet

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: my-daemonset
  namespace: kube-system
spec:
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      nodeSelector:
        node-
role.kubernetes.io/master: ""
      tolerations:
        - key: "node-
role.kubernetes.io/master"
          effect: NoSchedule
      containers:
        - name: my-container
          image: nginx:1.14.2

```

Job

```

apiVersion: batch/v1
kind: Job
metadata:
  name: my-job
spec:
  completions: 5
  parallelism: 2
  backoffLimit: 6
  template:
    spec:
      containers:
        - name: my-job-container
          image: busybox
          command: ["/bin/sh", "-c",
"echo Hello, Kubernetes! && sleep 30"]
          restartPolicy: OnFailure

```

CronJob

```

apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: my-cronjob
spec:
  schedule: "*/1 * * * *"
  concurrencyPolicy: Forbid
  failedJobsHistoryLimit: 1
  successfulJobsHistoryLimit: 3
  suspend: false
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: my-cronjob-
container
              image: busybox
              command:
                - /bin/sh
                - -c
                - date; echo Hello
from the Kubernetes cron job
              restartPolicy: OnFailure

```

StatefulSet

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-statefulset

```



```
spec:
  serviceName: "my-service"
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
          volumeClaimTemplates:
            - metadata:
                name: www
              spec:
                accessModes: ["ReadWriteOnce"]
                resources:
                  requests:
                    storage: 1Gi
```

Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    kubernetes.io/ingress.class: "nginx"
    nginx.ingress.kubernetes.io/rewrite-target: /$2
spec:
  rules:
    - host: example.com
      http:
        paths:
          - path: /path/(.*)
            pathType: Prefix
            backend:
              service:
                name: my-service
                port:
                  number: 80
```

HorizontalPodAutoscaler

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoscaler
metadata:
  name: my-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-deployment
  minReplicas: 1
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 50
    - type: Pods
      pods:
        metric:
          name: packets-per-second
        target:
          type: AverageValue
          averageValue: 1k
```

VerticalPodAutoscaler

Requires additional installation of Vertical Pod Autoscaler:

```
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: my-vpa
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: my-deployment
  updatePolicy:
    updateMode: "Auto"
  resourcePolicy:
    containerPolicies:
      - containerName: '*'
        minAllowed:
          cpu: 250m
          memory: 64Mi
        maxAllowed:
          cpu: 2
          memory: 4Gi
```

NetworkPolicy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-network-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: frontend
      ports:
        - protocol: TCP
          port: 6379
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
            except:
              - 10.0.0.0/28
      ports:
        - protocol: TCP
          port: 5978
```

ServiceAccount

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-service-account
  namespace: my-namespace
secrets:
  - name: my-secret
imagePullSecrets:
  - name: regcred
```

Endpoints

```
apiVersion: v1
kind: Endpoints
metadata:
  name: my-endpoints
subsets:
  - addresses:
      - ip: 192.168.1.1
        nodeName: worker1
    ports:
      - port: 80
        name: http
```

ResourceQuota

```
apiVersion: v1
```

```
kind: ResourceQuota
metadata:
  name: my-quota
spec:
  hard:
    pods: "10"
    requests.cpu: "4"
    requests.memory: 6Gi
    limits.cpu: "10"
    limits.memory: 10Gi
    configmaps: "10"
    secrets: "10"
    services: "5"
    services.loadbalancers: "1"
```

LimitRange

```
apiVersion: v1
kind: LimitRange
metadata:
  name: my-limitrange
spec:
  limits:
    - type: Pod
      max:
        cpu: "2"
        memory: 1Gi
      min:
        cpu: 200m
        memory: 6Mi
    - type: Container
      default:
        cpu: 500m
        memory: 512Mi
      defaultRequest:
        cpu: 100m
        memory: 128Mi
```

Roles and RoleBindings

```
# Role
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: my-namespace
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]

# RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: my-namespace
subjects:
- kind: User
  name: my-user
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

ClusterRoles and ClusterRoleBindings

```
# ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: secret-reader
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]

# ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
```

```
metadata:
  name: read-secrets-global
subjects:
- kind: User
  name: my-user
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io
```

CustomResourceDefinition

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: crontabs.stable.example.com
spec:
  group: stable.example.com
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                cronSpec:
                  type: string
                image:
                  type: string
                replicas:
                  type: integer

          subresources:
            status: {}
  scope: Namespaced
  names:
    plural: crontabs
    singular: crontab
    kind: CronTab
    shortNames:
      - ct
```

StorageClass

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
  zones: us-west-2a, us-west-2b
reclaimPolicy: Retain
allowVolumeExpansion: true
mountOptions:
- debug
volumeBindingMode: WaitForFirstConsumer
```

PodDisruptionBudget

```
apiVersion: policy/v1beta1
kind: PodDisruptionBudget
metadata:
  name: my-pdb
spec:
  minAvailable: 2
  selector:
    matchLabels:
      app: my-app
```

Use `kubectl apply -f <filename>.yaml` to deploy them to your Kubernetes cluster.

- **Pod:** `kubectl run my-pod --image=nginx --port=80 --restart=Never --dry-run=client -o yaml > pod.yaml`
- **Service:** `kubectl expose deployment my-deployment --type=LoadBalancer --`

```
port=80 --target-port=8080 --name=my-service
```

- **ConfigMap:** `kubectl create configmap my-config --from-file=config.properties`
- **Secret:** `kubectl create secret generic my-secret --from-literal=username=user --from-literal=password=pass`
- **Namespace:** `kubectl create namespace my-namespace`
- **Deployment:** `kubectl create deployment my-deployment --image=nginx --replicas=3 --dry-run=client -o yaml > deploy.yaml`
- **HorizontalPodAutoscaler:** `kubectl autoscale deployment my-deployment --min=1 --max=10 --cpu-percent=50`

Important Kubernetes Commands:

- **kubectl get pods:** Lists all pods in the current namespace.
 - Syntax: `kubectl get pods [-n <namespace>] [-o <output_format>]`
 - Example: `kubectl get pods -n default -o wide`
- **kubectl get nodes:** Shows all nodes in the cluster.
 - Syntax: `kubectl get nodes [-o <output_format>]`
 - Example: `kubectl get nodes -o json`
- **kubectl get services:** Lists all services in the current namespace.
 - Syntax: `kubectl get services [-n <namespace>] [-o <output_format>]`
 - Example: `kubectl get services -n kube-system`
- **kubectl describe pod :** Provides detailed information about a specific pod.
 - Syntax: `kubectl describe pod <pod-name> [-n <namespace>]`
 - Example: `kubectl describe pod my-pod -n my-namespace`
- **kubectl logs :** Retrieves logs from a container in a pod.
 - Syntax: `kubectl logs <pod-name> [-c <container-name>] [--previous] [-f]`
 - Example: `kubectl logs my-pod -c my-container --previous`
- **kubectl exec -it -- /bin/bash:** Opens an interactive shell into a container within a pod.
 - Syntax: `kubectl exec -it <pod-name> [-c <container-name>] -- <command>`
 - Example: `kubectl exec -it my-pod -c main-container -- /bin/bash`
- **kubectl apply -f .yaml:** Applies a configuration to a resource by filename or stdin.
 - Syntax: `kubectl apply -f <filename>.yaml [-n <namespace>]`
 - Example: `kubectl apply -f deployment.yaml`
- **kubectl delete pod :** Deletes a pod.
 - Syntax: `kubectl delete pod <pod-name> [-n <namespace>]`

- Example: `kubectl delete pod my-pod`

- **kubectl scale --replicas=3 deployment/:** Scales the number of pods for a deployment.
 - Syntax: `kubectl scale --replicas=<number> deployment/<deployment-name> [-n <namespace>]`
 - Example: `kubectl scale --replicas=3 deployment/my-app`
- **kubectl rollout status deployment/:** Checks the status of a deployment rollout.
 - Syntax: `kubectl rollout status deployment/<deployment-name> [-n <namespace>]`
 - Example: `kubectl rollout status deployment/my-deployment`
- **kubectl rollout undo deployment/:** Rolls back to the previous deployment revision.
 - Syntax: `kubectl rollout undo deployment/<deployment-name> [-n <namespace>]`
 - Example: `kubectl rollout undo deployment/my-deployment`
- **kubectl create deployment --image=:** Creates a new deployment with the specified image.
 - Syntax: `kubectl create deployment <deployment-name> --image=<image-name> [-n <namespace>]`
 - Example: `kubectl create deployment nginx --image=nginx`
- **kubectl get deployments:** Lists all deployments in the current namespace.
 - Syntax: `kubectl get deployments [-n <namespace>] [-o <output_format>]`
 - Example: `kubectl get deployments -o yaml`
- **kubectl port-forward ::** Forwards traffic from a local port to a port on the pod.
 - Syntax: `kubectl port-forward <pod-name> <local-port>:<pod-port> [-n <namespace>]`
 - Example: `kubectl port-forward my-pod 8080:80`
- **kubectl label nodes =:** Adds or updates a label on a node.
 - Syntax: `kubectl label nodes <node-name> <key>=<value> [--overwrite]`
 - Example: `kubectl label nodes worker1 disktype=ssd`
- **kubectl taint nodes =::** Adds a taint on a node, which can repel pods unless they tolerate the taint.
 - Syntax: `kubectl taint nodes <node-name> <key>=<value>:<effect> [--overwrite]`
 - Example: `kubectl taint nodes worker2 apptype=legacy:NoSchedule`
- **kubectl get events:** Shows all events in the current namespace.

- o Syntax: `kubectl get events [-n <namespace>] [-o <output_format>]`
- o Example: `kubectl get events -n my-namespace --sort-by='.lastTimestamp'`
- **kubectl config view:** Displays current kubeconfig settings.
 - o Syntax: `kubectl config view [--minify] [--flatten]`
 - o Example: `kubectl config view --minify`
- **kubectl cluster-info:** Displays endpoint information about the master and services in the cluster.
 - o Syntax: `kubectl cluster-info`
 - o Example: `kubectl cluster-info`

How do you mount a ConfigMap as an environment variable or volume in a Pod?

For environment variables:

```
env:
- name: SPECIAL_LEVEL_KEY
  valueFrom:
    configMapKeyRef:
      name: special-config
      key: special.how
```

For volumes:

```
volumes:
- name: config-volume
  configMap:
    name: special-config
volumeMounts:
- mountPath: /etc/config
  name: config-volume
```

How would you securely use Secrets in a Pod?

Mount Secrets as files in a volume for minimal exposure or use them as environment variables. For file mounts:

```
volumes:
- name: secret-volume
  secret:
    secretName: mysecret
volumeMounts:
- name: secret-volume
  readOnly: true
  mountPath: "/etc/secrets"
```

For environment variables:

```
env:
- name: SECRET_USERNAME
  valueFrom:
    secretKeyRef:
      name: mysecret
      key: username
```

How can you schedule Pods on specific nodes?

Use nodeSelector in the pod spec to match node labels:

```
nodeSelector:
  disktype: ssd
```

Or use nodeAffinity for more complex rules. Taints and tolerations can also be used to repel or attract pods to nodes.

Describe how you would configure an Ingress to route traffic to different services.

Define rules in the Ingress resource:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
spec:
  rules:
  - host: example.com
    http:
```

```
paths:
- path: /api
  pathType: Prefix
  backend:
    service:
      name: api-service
      port:
        number: 80
- path: /
  pathType: Prefix
  backend:
    service:
      name: web-service
      port:
        number: 80
```

How do you implement a NetworkPolicy to restrict pod communication?

Define a NetworkPolicy with selectors and rules for ingress/egress:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-from-namespace
spec:
  podSelector:
    matchLabels:
      role: frontend
  policyTypes:
  - Ingress
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          environment: production
    ports:
    - protocol: TCP
      port: 80
```

How can you bind a ServiceAccount to a Role or ClusterRole?

Use RoleBindings or ClusterRoleBindings:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: ServiceAccount
  name: my-service-account
  namespace: default
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

How would you set up ResourceQuotas to prevent a namespace from using too many resources?

Define a ResourceQuota in the namespace:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
spec:
  hard:
    pods: "4"
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```

Update kubeconfig for EKS:

```
aws eks update-kubeconfig --name my-cluster --region ap-south-1
```

The script sets up port forwarding for Prometheus to access it locally.

```
kubectl port-forward $(kubectl get pods -l
app=prometheus -o
jsonpath='{.items[0].metadata.name}')
9090:9090 > /dev/null 2>&1 &
```

mysql-secret.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: mysql-secret
type: Opaque
data:
  MYSQL_ROOT_PASSWORD: cGFzc3dvcmQ= #
base64 encoded value of "password"
```

backend-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: backend-service
spec:
  selector:
    app: backend
  ports:
    - protocol: TCP
      port: 3000
      targetPort: 3000
  type: LoadBalancer
```

backenddeployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: backend
  template:
    metadata:
      labels:
        app: backend
    spec:
      initContainers:
        - name: init-mysql
          image: mysql:8.0
          env:
            - name: MYSQL_ROOT_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-
secret
                  key:
MYSQL_ROOT_PASSWORD
- name: DB_HOST
  value: "${db_host}"
- name: DB_PORT
  value: "${db_port}"
volumeMounts:
- name: init-sql
  mountPath: /docker-
entrypoint-initdb.d
  command: [ "sh", "-c", "mysql
-h ${db_host} -P ${db_port} -u admin -
p${MYSQL_ROOT_PASSWORD} < /docker-entrypoint-
initdb.d/init.sql" ]
  containers:
    - name: backend
      image:
jeevan2001/backend:latest
      env:
        - name: DB_HOST
          value: "${db_host}"
        - name: DB_PORT
          value: "${db_port}"
        - name: MYSQL_ROOT_PASSWORD
          valueFrom:
```

```
secretKeyRef:
  name: mysql-
```

```
secret
```

```
key:
```

```
MYSQL_ROOT_PASSWORD
```

```
ports:
```

```
- containerPort: 3000
```

```
volumes:
```

```
- name: init-sql
```

```
configMap:
```

```
name: init-sql-config
```

Get the Backend LoadBalancer DNS

```
export BACKEND_LOADBALANCER_DNS=$(kubectl get
service backend-service -o
jsonpath='{.status.loadBalancer.ingress[0].ho
stname}')
```

frontendservice.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
spec:
  selector:
    app: frontend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

frontenddeployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: frontend
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
        - name: frontend
          image:
jeevan2001/frontend:latest
          ports:
            - containerPort: 80
          imagePullPolicy: Always
```

hpa-backend.yaml

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: hpa-backend
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: backend-deployment
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

cluster-autoscaler.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: cluster-autoscaler
  namespace: kube-system
  labels:
    app: cluster-autoscaler
spec:
  replicas: 1
  selector:
```

```

    matchLabels:
      app: cluster-autoscaler
  template:
    metadata:
      labels:
        app: cluster-autoscaler
    spec:
      containers:
        - name: cluster-autoscaler
          image:
            k8s.gcr.io/autoscaling/cluster-autoscaler:v1.20.0
          command:
            - ./cluster-autoscaler
            - --v=4
            - --stderrthreshold=info
            - --cloud-provider=aws
            - --skip-nodes-with-local-storage=false
            - --expander=least-waste
            - --nodes=1:10:my-node-group
          env:
            - name: AWS_REGION
              value: ap-south-1
          resources:
            limits:
              cpu: 100m
              memory: 300Mi
            requests:
              cpu: 100m
              memory: 300Mi
          volumeMounts:
            - name: ssl-certs
              mountPath:
                /etc/ssl/certs/ca-certificates.crt
              readOnly: true
          volumes:
            - name: ssl-certs
              hostPath:
                path: /etc/ssl/certs/ca-certificates.crt

```

cluster-autoscaler-policy.json

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "autoscaling:DescribeAutoScalingGroups",
        "autoscaling:DescribeAutoScalingInstances",
        "autoscaling:DescribeLaunchConfigurations",
        "autoscaling:DescribeTags",
        "autoscaling:SetDesiredCapacity",
        "autoscaling:TerminateInstanceInAutoScalingGroup",
        "ec2:DescribeLaunchTemplateVersions"
      ],
      "Resource": "*",
      "Effect": "Allow"
    }
  ]
}

```

Terraform

AWS Provider

```

provider "aws" {
  region = "ap-south-1"
}

```

Kubernetes Provider

```

provider "kubernetes" {

```

```

  host =
    aws_eks_cluster.my_cluster.endpoint
  cluster_ca_certificate =
    base64decode(aws_eks_cluster.my_cluster.certificate_authority[0].data)
  token =
    data.aws_eks_cluster_auth.my_cluster.token
}

```

Data Sources

aws_eks_cluster_auth

```

data "aws_eks_cluster_auth" "my_cluster" {
  name = aws_eks_cluster.my_cluster.name
}

```

aws_availability_zones

```

data "aws_availability_zones" "available" {}

```

Network Resources

aws_vpc

```

resource "aws_vpc" "eks_vpc" {
  cidr_block = "10.0.0.0/16"
}

```

aws_subnet

```

resource "aws_subnet" "eks_public_subnet" {
  count = 3
  vpc_id =
    aws_vpc.eks_vpc.id
  cidr_block =
    cidrsubnet(aws_vpc.eks_vpc.cidr_block, 8, count.index)
  availability_zone =
    element(data.aws_availability_zones.available.names, count.index)
  map_public_ip_on_launch = true
}

```

aws_subnet (Private)

```

resource "aws_subnet" "eks_private_subnet" {
  count = 3
  vpc_id =
    aws_vpc.eks_vpc.id
  cidr_block =
    cidrsubnet(aws_vpc.eks_vpc.cidr_block, 8, count.index + 3)
  availability_zone =
    element(data.aws_availability_zones.available.names, count.index)
  map_public_ip_on_launch = false
}

```

aws_internet_gateway

```

resource "aws_internet_gateway" "eks_igw" {
  vpc_id = aws_vpc.eks_vpc.id
}

```

aws_route_table

```

resource "aws_route_table"
"eks_public_route_table" {
  vpc_id = aws_vpc.eks_vpc.id

  route {
    cidr_block = "0.0.0.0/0"
    gateway_id =
      aws_internet_gateway.eks_igw.id
  }
}

```

aws_route_table_association

```

resource "aws_route_table_association"
"eks_public_route_table_association" {
  count = 3
  subnet_id =
    element(aws_subnet.eks_public_subnet[*].id, count.index)
  route_table_id =
    aws_route_table.eks_public_route_table.id
}

```

aws_nat_gateway

```

resource "aws_nat_gateway" "eks_nat_gateway" {
  count = 3

```

```

        allocation_id =
aws_eip.nat_eip[count.index].id
        subnet_id      =
element(aws_subnet.eks_public_subnet[*].id,
count.index)
}

```

aws_eip

```

resource "aws_eip" "nat_eip" {
    count    = 3
    domain   = "vpc"
}

```

aws_route_table (Private)

```

resource "aws_route_table"
"eks_private_route_table" {
    vpc_id = aws_vpc.eks_vpc.id

    route {
        cidr_block      = "0.0.0.0/0"
        nat_gateway_id =
element(aws_nat_gateway.eks_nat_gateway[*].id
, 0)
    }
}

```

aws_route_table_association (Private)

```

resource "aws_route_table_association"
"eks_private_route_table_association" {
    count    = 3
    subnet_id =
element(aws_subnet.eks_private_subnet[*].id,
count.index)
    route_table_id =
aws_route_table.eks_private_route_table.id
}

```

Security

aws_security_group

```

resource "aws_security_group"
"eks_security_group" {
    vpc_id = aws_vpc.eks_vpc.id

    egress {
        from_port    = 0
        to_port      = 0
        protocol     = "-1"
        cidr_blocks  = ["0.0.0.0/0"]
    }

    ingress {
        from_port    = 3306
        to_port      = 3306
        protocol     = "tcp"
        cidr_blocks  = ["10.0.0.0/16"]
    }
}

```

Database

aws_db_instance

```

resource "aws_db_instance" "mydb" {
    allocated_storage    = 20
    storage_type         = "gp2"
    engine               = "mysql"
    engine_version       = "8.0"
    instance_class       = "db.t3.micro"
    db_name              = "mydatabase"
    username             = "admin"
    password             = "password"
    db_subnet_group_name =
aws_db_subnet_group.mydb_subnet_group.name
    vpc_security_group_ids =
[aws_security_group.rds_security_group.id]
    skip_final_snapshot = true
}

```

aws_db_subnet_group

```

resource "aws_db_subnet_group"
"mydb_subnet_group" {
    name = "mydb-subnet-group"
}

```

```

        subnet_ids =
aws_subnet.eks_private_subnet[*].id
}

```

IAM

aws_iam_role

```

resource "aws_iam_role" "eks_cluster_role" {
    name = "eks-cluster-role"

    assume_role_policy = jsonencode({
        Version = "2012-10-17"
        Statement = [
            {
                Effect = "Allow"
                Principal = {
                    Service =
"eks.amazonaws.com"
                }
                Action = "sts:AssumeRole"
            },
        ]
    })
}

```

aws_iam_role_policy_attachment

```

resource "aws_iam_role_policy_attachment"
"eks_cluster_role_attachment" {
    role      =
aws_iam_role.eks_cluster_role.name
    policy_arn =
"arn:aws:iam::aws:policy/AmazonEKSClusterPolicy"
}

```

EKS

aws_eks_cluster

```

resource "aws_eks_cluster" "my_cluster" {
    name     = "my-cluster"
    role_arn =
aws_iam_role.eks_cluster_role.arn

    vpc_config {
        subnet_ids =
aws_subnet.eks_public_subnet[*].id
        security_group_ids =
[aws_security_group.eks_security_group.id]
    }
}

```

aws_eks_node_group

```

resource "aws_eks_node_group" "my_node_group"
{
    cluster_name =
aws_eks_cluster.my_cluster.name
    node_group_name = "my-node-group"
    node_role_arn =
aws_iam_role.eks_node_role.arn
    subnet_ids =
aws_subnet.eks_private_subnet[*].id

    scaling_config {
        desired_size = 5
        max_size     = 7
        min_size     = 3
    }

    instance_types = ["t3.small"]

    remote_access {
        ec2_ssh_key = "my-key"
    }

    tags = {
        Name = "eks-node-group"
    }
}

```

Local Resources and Data

local_file

```
resource "local_file"
"website_content_configmap" {
  content =
data.template_file.website_content_configmap.
rendered
  filename = "${path.module}/website-
content-configmap.yaml"
}
```

data.template_file

```
data "template_file"
"website_content_configmap" {
  template = file("${path.module}/website-
content-configmap.tpl.yaml")
  vars = {
    db_host =
aws_db_instance.mydb.endpoint
  }
}
```

kubernetes_config_map

```
resource "kubernetes_config_map"
"init_sql_config" {
  metadata {
    name = "init-sql-config"
  }
  data = {
    "init.sql" =
file("${path.module}/init.sql")
  }
}
```

VPC

```
resource "aws_vpc" "eks_vpc" {
  cidr_block = "10.0.0.0/16"
}

resource "aws_subnet" "eks_public_subnet" {
  count                = 3
  vpc_id              =
aws_vpc.eks_vpc.id
  cidr_block          =
cidrsubnet(aws_vpc.eks_vpc.cidr_block, 8,
count.index)
  availability_zone    =
element(data.aws_availability_zones.available
.names, count.index)
  map_public_ip_on_launch = true
}

resource "aws_subnet" "eks_private_subnet" {
  count                = 3
  vpc_id              =
aws_vpc.eks_vpc.id
  cidr_block          =
cidrsubnet(aws_vpc.eks_vpc.cidr_block, 8,
count.index + 3)
  availability_zone    =
element(data.aws_availability_zones.available
.names, count.index)
}
```

Security Groups

AWS Security Group:

```
resource "aws_security_group"
"eks_security_group" {
  vpc_id = aws_vpc.eks_vpc.id

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
  }
}
```

```
cidr_blocks = ["0.0.0.0/0"]
}
```

Kubernetes Network Policy:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-web
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: web
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: frontend
      ports:
        - protocol: TCP
          port: 80
```

EKS Cluster

EKS Cluster:

```
resource "aws_eks_cluster" "my_cluster" {
  name     = "my-cluster"
  role_arn =
aws_iam_role.eks_cluster_role.arn

  vpc_config {
    subnet_ids =
[aws_subnet.eks_public_subnet.*.id]
  }
}
```

IAM Role for EKS Cluster:

```
resource "aws_iam_role" "eks_cluster_role" {
  name = "eks-cluster-role"

  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [
      {
        Effect = "Allow"
        Principal = {
          Service =
"eks.amazonaws.com"
        }
        Action = "sts:AssumeRole"
      }
    ]
  })
}

resource "aws_iam_role_policy_attachment"
"eks_cluster_policy" {
  role =
aws_iam_role.eks_cluster_role.name
  policy_arn =
"arn:aws:iam::aws:policy/AmazonEKSClusterPolicy"
}
```

AWS & Kubernetes Integration with Terraform

```
provider "aws" {
  region = "ap-south-1"
}

provider "kubernetes" {
  host =
aws_eks_cluster.my_cluster.endpoint
  cluster_ca_certificate =
base64decode(aws_eks_cluster.my_cluster.certi
ficate_authority[0].data)
  token =
data.aws_eks_cluster_auth.my_cluster.token
}
```



```
resource "aws_eks_cluster" "my_cluster" {
  name      = "my-cluster"
  role_arn =
aws_iam_role.eks_cluster_role.arn

  vpc_config {
    subnet_ids =
[aws_subnet.eks_public_subnet.*.id]
  }
}
```

Code Example:

ConfigMap:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: db-config
data:
  DB_HOST: mydb.example.com
  DB_PORT: "3306"
```

Secret:

```
apiVersion: v1
kind: Secret
metadata:
  name: db-secret
type: Opaque
data:
  DB_PASSWORD: cGFzc3dvcmQ= # base64 encoded
password
```

Using ConfigMap and Secret in a Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
  - name: my-app-container
    image: my-app-image
    env:
    - name: DB_HOST
      valueFrom:
        configMapKeyRef:
          name: db-config
          key: DB_HOST
    - name: DB_PORT
      valueFrom:
        configMapKeyRef:
          name: db-config
          key: DB_PORT
    - name: DB_PASSWORD
      valueFrom:
        secretKeyRef:
          name: db-secret
          key: DB_PASSWORD
```

Autoscaling using Kubernetes and AWS

AWS Auto Scaling Group:

```
resource "aws_autoscaling_group" "example" {
  launch_configuration =
aws_launch_configuration.example.id
  min_size             = 1
  max_size             = 5
  desired_capacity     = 2
  vpc_zone_identifier =
[aws_subnet.eks_public_subnet.*.id]
}
```

Kubernetes HPA:

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: my-app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  minReplicas: 1
```

maxReplicas: 10

targetCPUUtilizationPercentage: 50

- **Pods:** The smallest and simplest Kubernetes object. A Pod represents a single instance of a running process in your cluster.
- **ReplicaSets:** Ensures a specified number of pod replicas are running at any given time.
- **Deployments:** Provides declarative updates for Pods and ReplicaSets.
- **Services:** An abstraction which defines a logical set of Pods and a policy by which to access them - like load-balancers.
- **ConfigMaps:** Used to store configuration data in key-value pairs which can be consumed by pods.
- **Secrets:** Manages sensitive information, like passwords, OAuth tokens, and ssh keys, which can be referenced in pod definitions.
- **PersistentVolumes (PV):** A piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.
- **PersistentVolumeClaims (PVC):** Requests storage resources defined by a PersistentVolume.
- **Namespaces:** Provides a scope for names. Resources like Pods, Services, and Deployments can be isolated within namespaces.
- **Nodes:** A worker machine in Kubernetes, either virtual or physical, where containers will be launched by Kubernetes.
- **DaemonSets:** Ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected.
- **Jobs:** Creates one or more Pods and ensures that a specified number of them successfully terminate. Good for batch processes.
- **CronJobs:** Manages time-based Jobs, similar to cron in Unix-like systems.
- **StatefulSets:** Manages the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods.
- **Ingress:** Manages external access to the services in a cluster, typically HTTP.
- **HorizontalPodAutoscaler:** Scales a Deployment, ReplicaSet, or ReplicationController based on observed CPU utilization or other select metrics.
- **VerticalPodAutoscaler:** Automatically adjusts the compute resources of pods based on usage.
- **NetworkPolicies:** Specifies how groups of pods are allowed to communicate with each other and other network endpoints.

- **ServiceAccounts**: Provides an identity for processes that run in a Pod, which can be used for authenticating to the API server.
- **Endpoints**: Exposes the IP addresses of a service's backing pods.
- **ResourceQuotas**: Provides constraints that limit aggregate resource consumption per namespace.
- **LimitRanges**: Constrains resource allocations (to Pods or Containers) in a namespace.
- **Roles and RoleBindings (for RBAC - Role-Based Access Control)**: Define permissions for users or service accounts within a namespace.
- **ClusterRoles and ClusterRoleBindings**: Similar to Roles but cluster-wide, not namespace-specific.
- **CustomResourceDefinitions (CRDs)**: Allows users to create new types of resources without adding another API server.
- **StorageClasses**: Describes different classes or profiles of storage in the cluster.
- **PodDisruptionBudgets**: Ensures that a specified number of pods are available even during voluntary disruptions like node drains or upgrades.

----- Priority Order of Learning Kubernetes Resources (Quickie) -----

Priority 1: Must-Know Kubernetes Resources for Interviews -----

```
Pod
Deployment
Service
ConfigMap
Secret
PersistentVolume
PersistentVolumeClaim
Namespace
StatefulSet
Ingress
HorizontalPodAutoscaler
```

----- Priority 2: Nice-to-Know Resources (Learn if You Have Time) -----

```
Replicaset
DaemonSet
Job and CronJob
NetworkPolicy
ServiceAccount
ResourceQuota
LimitRange
```

----- Priority 3: Skip for Now (Unless Specialized) -----

```
VerticalPodAutoscaler
PodDisruptionBudget
CustomResourceDefinition
StorageClass
Endpoints
Roles
RoleBindings
ClusterRoles
ClusterRoleBindings
```

----- Priority 1: Must-Know Kubernetes Resources for Interviews -----

Pod

The smallest and simplest Kubernetes object. A Pod represents a single instance of a running process in your cluster.

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-pod
  labels:
    app: my-app
spec:
  containers:
    - name: app-container
      image: nginx:latest
      ports:
        - containerPort: 80
      resources:
        requests:
          cpu: "100m"
          memory: "128Mi"
        limits:
          cpu: "500m"
          memory: "256Mi"
```

Deployment

Provides declarative updates for Pods and ReplicaSets.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-deployment
  labels:
    app: my-app
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          resources:
            requests:
              memory: "256Mi"
              cpu: "200m"
            limits:
              memory: "512Mi"
              cpu: "500m"
          livenessProbe:
            httpGet:
              path: /health
              port: 80
            initialDelaySeconds: 30
            periodSeconds: 10
          readinessProbe:
            httpGet:
              path: /ready
              port: 80
            initialDelaySeconds: 5
            periodSeconds: 5
          env:
            - name: ENVIRONMENT
              value: "production"
```

Service

An abstraction which defines a logical set of Pods and a policy by which to access them - like loadbalancers.

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
```

```

selector:
  app: my-app
ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
    name: http
  type: LoadBalancer

```

ConfigMap

Used to store configuration data in key-value pairs which can be consumed by pods.

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: my-config
data:
  app.env: "production"
  config.file: |
    key1=value1
    key2=value2

```

Secret

Manages sensitive information, like passwords, OAuth tokens, and ssh keys, which can be referenced in pod definitions.

```

apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  username: YWRtaW4= # "admin"
  password: UEA1NXcwcmQ= # "P@55w0rd"

```

PersistentVolume

A piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0001
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  storageClassName: standard
  nfs:
    server: nfs-server.example.com
    path: "/exports"

```

PersistentVolumeClaim

Requests storage resources defined by a PersistentVolume.

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: standard

```

Namespace

Provides a scope for names. Resources like Pods, Services, and Deployments can be isolated within namespaces.

```

apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace
  labels:
    environment: production

```

StatefulSet

Manages the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods.

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-statefulset
spec:
  serviceName: my-service
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          volumeMounts:
            - name: www
              mountPath: "/usr/share/nginx/html"
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes:
          - ReadWriteOnce
        resources:
          requests:
            storage: 1Gi

```

Ingress

Manages external access to the services in a cluster, typically HTTP.

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
    - host: example.com
      http:
        paths:
          - path: /app
            pathType: Prefix
            backend:
              service:
                name: my-service
                port:
                  number: 80

```

HorizontalPodAutoscaler

Scales a Deployment, ReplicaSet, or ReplicationController based on observed CPU utilization or other select metrics.

```

apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: my-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-deployment
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70

```

Priority 2: Nice-to-Know Resources (Learn if You Have Time)

ReplicaSet

Ensures a specified number of pod replicas are running at any given time.

```

apiVersion: apps/v1
kind: ReplicaSet

```

```

metadata:
  name: my-replicaset
  labels:
    app: my-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:1.14.2
          ports:
            - containerPort: 80

```

DaemonSet

Ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected.

```

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: my-daemonset
spec:
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      nodeSelector:
        kubernetes.io/role: worker
      tolerations:
        - key: "node-role.kubernetes.io/control-plane"
          effect: "NoSchedule"
      containers:
        - name: my-container
          image: nginx:1.14.2

```

Job

Creates one or more Pods and ensures that a specified number of them successfully terminate. Good for batch processes.

```

apiVersion: batch/v1
kind: Job
metadata:
  name: my-job
spec:
  completions: 5
  parallelism: 2
  backoffLimit: 4
  template:
    spec:
      containers:
        - name: my-job-container
          image: busybox
          command: ["/bin/sh", "-c", "echo Hello, Kubernetes!"]
      restartPolicy: OnFailure

```

CronJob

Manages time-based Jobs, similar to cron in Unix-like systems.

```

apiVersion: batch/v1
kind: CronJob
metadata:
  name: my-cronjob
spec:
  schedule: "0 */1 * * *" # Every hour
  concurrencyPolicy: Forbid
  jobTemplate:
    spec:
      template:

```

```

spec:
  containers:
    - name: my-cronjob-container
      image: busybox
      command: ["/bin/sh", "-c", "echo Hello"]
  restartPolicy: OnFailure

```

NetworkPolicy

Specifies how groups of pods are allowed to communicate with each other and other network endpoints.

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-network-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: frontend
      ports:
        - protocol: TCP
          port: 6379
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
      ports:
        - protocol: TCP
          port: 3306

```

ServiceAccount

Provides an identity for processes that run in a Pod, which can be used for authenticating to the API server.

```

apiVersion: v1
kind: ServiceAccount
metadata:
  name: my-service-account
  namespace: devops-interview
imagePullSecrets:
  - name: regcred

```

ResourceQuota

Provides constraints that limit aggregate resource consumption per namespace.

```

apiVersion: v1
kind: ResourceQuota
metadata:
  name: my-quota
  namespace: devops-interview
spec:
  hard:
    pods: "10"
    requests.cpu: "4"
    requests.memory: "6Gi"
    limits.cpu: "10"
    limits.memory: "10Gi"

```

LimitRange

Constrains resource allocations (to Pods or Containers) in a namespace.

```

apiVersion: v1
kind: LimitRange
metadata:
  name: my-limitrange
  namespace: devops-interview
spec:
  limits:
    - type: Container
      max:
        cpu: "1"
        memory: "512Mi"
      min:
        cpu: "100m"

```

```

    memory: "64Mi"
  default:
    cpu: "500m"
    memory: "512Mi"
  defaultRequest:
    cpu: "200m"
    memory: "256Mi"

```

Priority 3: Skip for Now (Unless Specialized)

VerticalPodAutoscaler

Automatically adjusts the compute resources of pods based on usage.

```

apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: my-vpa
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: my-deployment
  updatePolicy:
    updateMode: "Auto"
  resourcePolicy:
    containerPolicies:
      - containerName: "*"
        minAllowed:
          cpu: "250m"
          memory: "128Mi"
        maxAllowed:
          cpu: "2"
          memory: "4Gi"

```

PodDisruptionBudget

Ensures that a specified number of pods are available even during voluntary disruptions like node drains or upgrades.

```

apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: my-pdb
spec:
  minAvailable: 2
  selector:
    matchLabels:
      app: my-app

```

CustomResourceDefinition

Allows users to create new types of resources without adding another API server.

```

apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: crontabs.stable.example.com
spec:
  group: stable.example.com
  scope: Namespaced
  names:
    plural: crontabs
    singular: crontab
    kind: CronTab
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                cronSpec:
                  type: string
                image:
                  type: string

```

StorageClass

Describes different classes or profiles of storage in the cluster.

```

apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:

```

```

  name: standard
  provisioner: kubernetes.io/aws-ebs
  parameters:
    type: gp2
  reclaimPolicy: Retain
  allowVolumeExpansion: true
  volumeBindingMode: WaitForFirstConsumer

```

Endpoints

Exposes the IP addresses of a service's backing pods.

```

apiVersion: v1
kind: Endpoints
metadata:
  name: my-endpoints
spec:
  subsets:
    - addresses:
        - ip: 192.168.1.1
      ports:
        - port: 80
          name: http

```

Roles

Define permissions for users or service accounts within a namespace.

```

apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: pod-reader
  namespace: devops-interview
rules:
  - apiGroups: [""]
    resources: ["pods"]
    verbs: ["get", "list", "watch"]

```

RoleBindings

Define permissions for users or service accounts within a namespace.

```

apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-pods
  namespace: devops-interview
subjects:
  - kind: User
    name: my-user
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io

```

ClusterRoles

Similar to Roles but cluster-wide, not namespace-specific.

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: secret-reader
rules:
  - apiGroups: [""]
    resources: ["secrets"]
    verbs: ["get", "list", "watch"]

```

ClusterRoleBindings

Similar to Roles but cluster-wide, not namespace-specific

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: read-secrets-global
subjects:
  - kind: User
    name: my-user
    apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: secret-reader
  apiGroup: rbac.authorization.k8s.io

```


Ansible

Program 1: Ansible Basics and Core Workflow

What is Ansible? (Core concepts, agentless, YAML, SSH)

Ansible is an open-source automation tool that uses an agentless architecture (no software installed on managed nodes), relies on SSH for communication, and uses YAML for configuration files like playbooks and inventory.

Ansible Inventory (Static vs. dynamic, host grouping)

The inventory defines the hosts Ansible manages. It can be static (a simple file) or dynamic (script-generated), with hosts organized into groups.

Ansible Ad-Hoc Commands (Basic usage, quick tasks)

Quick, one-line commands to perform tasks on hosts without writing a full playbook (e.g., `ansible all -m ping`).

Ansible Idempotence (Understanding the concept)

Ansible ensures tasks are idempotent, meaning running them multiple times produces the same result without unintended changes.

Step 1: Set Up a Static Inventory

Create a file named `hosts.ini` to define managed hosts and groups.

```
# File: hosts.ini
[webservers]
web1.example.com
web2.example.com

[dbservers]
db1.example.com
```

```
[all:vars]
ansible_user=admin
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

Explanation:

- `[webservers]` and `[dbservers]` are host groups.
- `ansible_user` and `ansible_ssh_private_key_file` are variables for SSH access, showcasing Ansible's agentless nature (uses SSH, no agents needed).

Step 2: Run an Ad-Hoc Command

Use an ad-hoc command to check host uptime.

```
ansible -i hosts.ini all -m command -a "uptime"
Output (example):
```

```
web1.example.com | SUCCESS | rc=0 >>
 14:35:23 up 5 days,  3:12,  1 user,  load
average: 0.10, 0.15, 0.20
web2.example.com | SUCCESS | rc=0 >>
 14:35:23 up 3 days, 10:45,  2 users,  load
average: 0.05, 0.08, 0.12
db1.example.com | SUCCESS | rc=0 >>
 14:35:23 up 7 days,  1:23,  1 user,  load
average: 0.25, 0.30, 0.35
```

Explanation:

- `-i hosts.ini`: Specifies the inventory file.
- `all`: Targets all hosts in the inventory.
- `-m command`: Uses the command module to execute uptime.

This demonstrates quick tasks without a playbook and Ansible's SSH-based communication.

Step 3: Demonstrate Idempotence

Run a simple idempotent ad-hoc command multiple times.

```
ansible -i hosts.ini webservers -m file -a
"path=/tmp/test.txt state=touch"
```

First Run Output:

```
web1.example.com | CHANGED => {"changed": true,
"path": "/tmp/test.txt"}
web2.example.com | CHANGED => {"changed": true,
"path": "/tmp/test.txt"}
```

Second Run Output:

```
web1.example.com | SUCCESS => {"changed": false,
"path": "/tmp/test.txt"}
web2.example.com | SUCCESS => {"changed": false,
"path": "/tmp/test.txt"}
```

Explanation:

- The file module with `state=touch` creates `/tmp/test.txt` if it doesn't exist.
- First run: File is created (`changed: true`).
- Second run: File already exists, so no change (`changed: false`), proving idempotence.

Step 4: Tie It Together

Ansible's core concepts are shown:

- Agentless: No software installed on `web1`, `web2`, or `db1`; SSH handles everything.
- YAML: Inventory uses a simple, readable format (though not strict YAML here, it's YAML-compatible).
- SSH: Connection relies on SSH keys defined in the inventory.

Key Takeaways for Notes:

- Inventory organizes hosts and groups for targeting.
- Ad-hoc commands are fast, playbook-free ways to manage systems.
- Idempotence ensures consistent results, a core Ansible principle.

Execution Command:

```
# Check connectivity
ansible -i hosts.ini all -m ping
# Run uptime command
ansible -i hosts.ini all -m command -a "uptime"
# Test idempotence
ansible -i hosts.ini webservers -m file -a
"path=/tmp/test.txt state=touch"
```

Program 2: Ansible Playbooks and Task Management

Topics Included:

- **Ansible Playbooks** (Structure, purpose, basic syntax)
 - Playbooks are YAML files defining a series of tasks to automate workflows.
- **Ansible Modules** (Common modules like `command`, `shell`, `copy`, `service`, `package`, `file`, `template`)
 - Modules are reusable units of work (e.g., `copy` for files, `service` for managing services).
- **Ansible Loops** (`loop`, basic iteration)
 - Loops allow repeating tasks over a list of items.
- **Ansible Conditionals** (`when`, basic operators)

- Conditionals control task execution based on conditions (e.g., OS type).
- **Ansible Tags** (Purpose, usage, running specific tasks)
 - Tags label tasks for selective execution.
- **Ansible Blocks** (Basic usage, grouping tasks)
 - Blocks group related tasks for better organization or error handling.
- **Ansible Command Module vs. Shell Module** (Differences, when to use which)
 - **command**: Runs simple commands without shell features.
 - **shell**: Runs commands with shell capabilities (e.g., pipes).

Step 1: Create a Playbook

Create a file named `setup_webserver.yml`.

```
# File: setup_webserver.yml
---
- name: Set up a basic web server
  hosts: webserver
  tasks:
    # Block for package installation
    - name: Install required packages
      block:
        - name: Install httpd and unzip
          ansible.builtin.package:
            name: "{{ item }}"
            state: present
          loop:
            - httpd
            - unzip
          tags: install

    # Task with conditional
    - name: Copy index.html to web server
      ansible.builtin.copy:
        src: ./files/index.html
        dest: /var/www/html/index.html
        mode: '0644'
        when: ansible_os_family == "RedHat"
        tags: configure

    # Task comparing command vs shell
    - name: Check httpd version with command
      module:
        ansible.builtin.command: httpd -v
        register: httpd_version_cmd
        tags: check

    - name: Check disk usage with shell module
      ansible.builtin.shell: df -h | grep /dev
      register: disk_usage
      tags: check

    # Service management
    - name: Ensure httpd is running
      ansible.builtin.service:
        name: httpd
        state: started
        enabled: yes
        tags: service
```

Explanation:

- **Playbook Structure**: Starts with `---`, defines a play targeting webserver.
- **Modules**: Uses package, copy, command, shell, and service.
- **Loops**: Installs multiple packages (httpd, unzip) with loop.

- **Conditionals**: Copies index.html only on RedHat-based systems.
- **Tags**: Labels tasks as install, configure, check, or service.
- **Blocks**: Groups package installation tasks.
- **Command vs. Shell**: command runs httpd -v (no shell needed); shell runs df -h | grep /dev (needs pipe).

Step 2: Prepare Supporting Files

Create a simple index.html file in a files/ directory.

```
<!-- File: files/index.html -->
<h1>Welcome to My Web Server</h1>
```

Step 3: Use an Inventory

Reuse the hosts.ini from Program 1 (assuming webserver group exists).

```
# File: hosts.ini
[webserver]
web1.example.com
web2.example.com
```

Step 4: Run the Playbook

Execute the full playbook:

```
ansible-playbook -i hosts.ini setup_webserver.yml
Run specific tagged tasks:
```

```
ansible-playbook -i hosts.ini setup_webserver.yml
--tags "install,configure"
Output (example):
```

```
TASK [Install httpd and unzip] *****
changed: [web1.example.com] => (item=httpd)
changed: [web1.example.com] => (item=unzip)
TASK [Copy index.html to web server] *****
changed: [web1.example.com]
TASK [Check httpd version with command module]
****
changed: [web1.example.com]
TASK [Check disk usage with shell module] ****
changed: [web1.example.com]
TASK [Ensure httpd is running] *****
changed: [web1.example.com]
```

Step 5: Verify Results

Check outputs stored in register: Add a debug task (optional) to see httpd_version_cmd and disk_usage:

```
- name: Debug outputs
  ansible.builtin.debug:
    var: httpd_version_cmd.stdout

- name: Debug disk usage
  ansible.builtin.debug:
    var: disk_usage.stdout
```

Rerun to see idempotence (most tasks show changed: false on second run).

Key Takeaways for Notes:

- **Playbooks**: Automate multi-step workflows in YAML.
- **Modules**: Building blocks for tasks (e.g., copy for files, service for daemons).
- **Loops**: Simplify repetitive tasks.
- **Conditionals**: Add logic to adapt to environments.
- **Tags**: Enable selective task execution.
- **Blocks**: Organize related tasks.
- **Command vs. Shell**: Use command for simple tasks, shell for complex shell features.

Execution Commands:

```
# Run full playbook
ansible-playbook -i hosts.ini setup_webserver.yml

# Run only installation and configuration
ansible-playbook -i hosts.ini setup_webserver.yml
--tags "install,configure"

# Run checks only
ansible-playbook -i hosts.ini setup_webserver.yml
--tags "check"
```

Program 3: Advanced Playbook Features and Reusability

Topics Included:

- **Ansible Roles** (Organization, reusability, basic structure)
 - Roles organize tasks, variables, and files into reusable units.
- **Ansible Variables** (Types, scope, usage)
 - Variables store dynamic data (e.g., package names) with different scopes (play, role, host).
- **Ansible Facts** (Purpose, usage, basic facts)
 - Facts are system details (e.g., OS, IP) gathered from managed nodes.
- **Ansible Handlers** (Purpose, usage, notify)
 - Handlers are tasks triggered by notify when changes occur (e.g., restart a service).
- **Ansible Templates** (Jinja2, basic usage)
 - Templates use Jinja2 to generate dynamic files (e.g., config files).

Step 1: Set Up a Role Structure

Create a role named webserver with the standard directory layout.

```
mkdir -p
roles/webserver/{tasks,handlers,templates,vars,files}
```

Explanation: Roles organize code into tasks/ (main logic), handlers/ (triggered tasks), templates/ (dynamic files), vars/ (variables), and files/ (static files).

Step 2: Define Role Components

Main Tasks (roles/webserver/tasks/main.yml):

```
---
- name: Install web server package
  ansible.builtin.package:
    name: "{{ web_package }}"
    state: present
  notify: Restart web service

- name: Copy static index.html
  ansible.builtin.copy:
    src: index.html
    dest: "{{ web_doc_root }}/index.html"
    mode: '0644'

- name: Generate httpd.conf from template
  ansible.builtin.template:
    src: httpd.conf.j2
    dest: /etc/httpd/conf/httpd.conf
    mode: '0644'
  notify: Restart web service
```

```
- name: Ensure web service is running
  ansible.builtin.service:
    name: "{{ web_service }}"
    state: started
    enabled: yes
```

Variables (roles/webserver/vars/main.yml):

```
---
web_package: httpd
web_service: httpd
web_doc_root: /var/www/html
```

Handlers

(roles/webserver/handlers/main.yml):

```
---
- name: Restart web service
  ansible.builtin.service:
    name: "{{ web_service }}"
    state: restarted
```

Template

(roles/webserver/templates/httpd.conf.j2):

```
Listen {{ ansible_default_ipv4.address }}:80
ServerName {{ ansible_hostname }}
DocumentRoot "{{ web_doc_root }}"
<Directory "{{ web_doc_root }}">
    AllowOverride All
    Require all granted
</Directory>
```

Static File

(roles/webserver/files/index.html):

```
<h1>Hello from {{ ansible_hostname }}!</h1>
Step 3: Create a Playbook to Use the Role
```

Create deploy_web.yml:

```
---
- name: Deploy web server using role
  hosts: webserver
  pre_tasks:
    - name: Gather facts
      ansible.builtin.setup:
    - name: Debug OS and IP
      ansible.builtin.debug:
        msg: "Running on {{
ansible_os_family }}" with IP {{
ansible_default_ipv4.address }}"
  roles:
    - webserver
```

Explanation:

- **Roles:** The webserver role is applied to webserver.
- **Variables:** web_package, web_service, etc., are defined in the role's vars/.
- **Facts:** ansible_os_family, ansible_hostname, and ansible_default_ipv4.address are used dynamically.
- **Handlers:** Notified when the package or config changes.
- **Templates:** httpd.conf.j2 uses Jinja2 to insert facts like IP and hostname.

Step 4: Use an Inventory

Reuse hosts.ini from previous programs:

```
# File: hosts.ini
[webserver]
web1.example.com
web2.example.com

[all:vars]
ansible_user=admin
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

Step 5: Run the Playbook

Execute the playbook:

```
ansible-playbook -i hosts.ini deploy_web.yml
Output (example):
```

```
TASK [Debug OS and IP] *****
ok: [web1.example.com] => {
  "msg": "Running on RedHat with IP
192.168.1.10"
}
```

```
TASK [webserver : Install web server package] ****
changed: [web1.example.com]
TASK [webserver : Copy static index.html] ****
changed: [web1.example.com]
TASK [webserver : Generate httpd.conf from
template] ****
changed: [web1.example.com]
TASK [webserver : Ensure web service is running]
****
changed: [web1.example.com]
HANDLER [webserver : Restart web service] ****
changed: [web1.example.com]
```

Step 6: Verify Results

On web1.example.com, check:

- `curl http://192.168.1.10`: Should show "Hello from web1!" (hostname from facts).
- `/etc/httpd/conf/httpd.conf`: Contains the IP and hostname from the template.

Key Takeaways for Notes:

- **Roles**: Modularize tasks for reusability (e.g., webserver role can be reused across projects).
- **Variables**: Define constants (e.g., web_package) in vars/ for flexibility.
- **Facts**: Automatically gather system info (e.g., ansible_hostname) for dynamic configs.
- **Handlers**: Trigger actions (e.g., service restart) only when needed.
- **Templates**: Use Jinja2 to create dynamic files based on facts and variables.

Execution Command:

```
ansible-playbook -i hosts.ini deploy_web.yml
```

Program 4: Security and Operational Control

Topics Included:

- **Ansible Vault** (Basic encryption, usage)
 - Vault encrypts sensitive data (e.g., passwords) in files.
- **Ansible Privilege Escalation** (become, become_user)
 - become escalates privileges (e.g., to root) for tasks requiring elevated access.
- **Ansible Check Mode (Dry Run)** (--check)
 - Check mode simulates tasks without making changes.
- **Ansible Best Practices** (Organization, security, readability)
 - Best practices include clear naming, modular structure, and secure handling of secrets.

Step 1: Encrypt Sensitive Data with Ansible Vault

Create an encrypted file `secrets.yml` for sensitive variables.

```
ansible-vault create secrets.yml
```

Enter a vault password (e.g., mypassword) when prompted, then add:

```
# File: secrets.yml
db_password: "securepass123"
```

Explanation: Vault encrypts `secrets.yml` to protect `db_password`.

Step 2: Create a Playbook with Security Features

Create secure_setup.yml:

```
---
- name: Securely set up a database server
  hosts: dbservers
  vars_files:
    - secrets.yml # Include encrypted
  variables:
    tasks:
      - name: Install MariaDB package
        ansible.builtin.package:
          name: mariadb-server
          state: present
          become: yes # Escalate privileges to
root
          become_user: root
          tags: install

      - name: Ensure MariaDB service is running
        ansible.builtin.service:
          name: mariadb
          state: started
          enabled: yes
          become: yes
          become_user: root
          tags: service

      - name: Set database root password
        ansible.builtin.shell: mysqladmin -u
root password "{{ db_password }}"
        when: ansible_os_family == "RedHat"
        become: yes
        become_user: root
        tags: configure
        no_log: true # Hide sensitive output
(best practice)
```

Explanation:

- **Vault**: `secrets.yml` provides `db_password`.
- **Privilege Escalation**: `become: yes` and `become_user: root` allow installing packages and managing services.
- **Check Mode**: Can be tested with `--check`.
- **Best Practices**:
 - Clear task names (e.g., "Install MariaDB package").
 - `no_log: true` hides sensitive data in logs.
 - Modular structure with tags (`install`, `service`, `configure`).

Step 3: Use an Inventory

Reuse or adapt `hosts.ini`:

```
# File: hosts.ini
[dbservers]
db1.example.com
[all:vars]
ansible_user=admin
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

Step 4: Run the Playbook

Dry Run (Check Mode):

```
ansible-playbook -i hosts.ini secure_setup.yml --
check --ask-vault-pass
```

Enter the vault password (mypassword) when prompted.

Output (example):

```
TASK [Install MariaDB package] *****
ok: [db1.example.com] => (skipped, in check mode)
TASK [Ensure MariaDB service is running] ****
ok: [db1.example.com] => (skipped, in check mode)
TASK [Set database root password] *****
ok: [db1.example.com] => (skipped, in check mode)
Full Execution:
ansible-playbook -i hosts.ini secure_setup.yml --
ask-vault-pass
Output (example):
TASK [Install MariaDB package] *****
```

```
changed: [db1.example.com]
TASK [Ensure MariaDB service is running] ****
changed: [db1.example.com]
TASK [Set database root password] *****
changed: [db1.example.com]
```

Step 5: Verify Results

On db1.example.com:

- Check if mariadb-server is installed (rpm -q mariadb-server).
- Verify MariaDB is running (systemctl status mariadb).
- Test the root password (mysql -u root -p with securepass123).

Key Takeaways for Notes:

- **Vault:** Encrypts sensitive data (e.g., db_password) for security.
- **Privilege Escalation:** become ensures tasks requiring root access succeed.
- **Check Mode:** --check previews changes without applying them.
- **Best Practices:**
 - Use descriptive names and tags.
 - Hide sensitive output with no_log.
 - Store secrets in Vault, not plaintext.

Execution Commands:

```
# Create/edit Vault file
ansible-vault edit secrets.yml --ask-vault-pass
# Dry run
ansible-playbook -i hosts.ini secure_setup.yml --check --ask-vault-pass
# Full run
ansible-playbook -i hosts.ini secure_setup.yml --ask-vault-pass
```

Program 5: Debugging and Validation

Topics Included:

- **Ansible Debugging** (Basic techniques, -v, debug module)
 - Debugging tools like verbose mode (-v) and the debug module help troubleshoot issues.
- **Ansible Check Mode (Dry Run)** (--check)
 - Check mode simulates playbook execution without applying changes.

Step 1: Create a Playbook for Debugging

Create debug_validate.yml:

```
---
- name: Debug and validate system setup
  hosts: webserver
  tasks:
    - name: Gather facts
      ansible.builtin.setup:
        tags: facts

    - name: Debug system OS and memory
      ansible.builtin.debug:
        msg: "OS: {{ ansible_os_family }},
Free Memory: {{ ansible_memfree_mb }} MB"
        tags: debug

    - name: Install httpd package
      ansible.builtin.package:
        name: httpd
        state: present
        register: install_result # Store task
        tags: install

    - name: Debug installation result
      ansible.builtin.debug:
```

```
var: install_result
when: install_result is defined
tags: debug
```

```
- name: Ensure httpd is running
  ansible.builtin.service:
    name: httpd
    state: started
    register: service_result
    tags: service

- name: Debug service status
  ansible.builtin.debug:
    msg: "Service changed: {{
service_result.changed }}, State: {{
service_result.state }}"
    when: service_result is defined
    tags: debug
```

Explanation:

- **Debugging:** Uses debug module to print facts (e.g., OS, memory) and task results.
- **Check Mode:** Can simulate package installation and service management.
- **Register:** Captures task outputs (install_result, service_result) for inspection.

Step 2: Use an Inventory

Reuse hosts.ini from previous programs:

```
# File: hosts.ini
[webserver]
web1.example.com

[all:vars]
ansible_user=admin
ansible_ssh_private_key_file=~/.ssh/id_rsa
```

Step 3: Run the Playbook with Debugging Verbose Mode (Basic):

```
ansible-playbook -i hosts.ini debug_validate.yml -v
```

Verbose Mode (Detailed):

```
ansible-playbook -i hosts.ini debug_validate.yml -vvv
```

Output (example with -v):

```
TASK [Debug system OS and memory] *****
ok: [web1.example.com] => {
  "msg": "OS: RedHat, Free Memory: 2048 MB"
}
TASK [Install httpd package] *****
changed: [web1.example.com] => {"changed": true,
"name": "httpd"}
TASK [Debug installation result] *****
ok: [web1.example.com] => {
  "install_result": {"changed": true,
"name": "httpd", "state": "present"}
}
```

Explanation:

- -v shows task outputs; -vvv adds detailed execution info (e.g., SSH commands).

Step 4: Run in Check Mode

Simulate execution:

```
ansible-playbook -i hosts.ini debug_validate.yml --check
```

Output (example):

```
TASK [Debug system OS and memory] *****
ok: [web1.example.com] => {
  "msg": "OS: RedHat, Free Memory: 2048 MB"
}
TASK [Install httpd package] *****
ok: [web1.example.com] => (skipped, in check mode)
TASK [Debug installation result] *****
skipping: [web1.example.com] # Skipped because
install_result isn't set in check mode
TASK [Ensure httpd is running] *****
ok: [web1.example.com] => (skipped, in check mode)
```

Explanation: Check mode runs debug tasks but skips changes (e.g., package install).

Step 5: Verify Debugging Output

Rerun with tags to focus on debugging:

```
ansible-playbook -i hosts.ini debug_validate.yml -  
-tags "debug" -v
```

Output (example):

```
TASK [Debug system OS and memory] *****  
ok: [web1.example.com] => {  
    "msg": "OS: RedHat, Free Memory: 2048 MB"  
}  
TASK [Debug installation result] *****  
ok: [web1.example.com] => {  
    "install_result": {"changed": false,  
    "name": "httpd"}  
}  
TASK [Debug service status] *****  
ok: [web1.example.com] => {  
    "msg": "Service changed: false, State:  
started"  
}
```

Key Takeaways for Notes:

- **Debugging:**
 - -v to -vvv: Increases verbosity for troubleshooting.
 - debug module: Prints variables, facts, or task results (e.g., `ansible_memfree_mb`).
- **Check Mode:** --check validates playbook logic without altering systems.
- Combine register with debug to inspect task outcomes.

Execution Commands:

```
# Run with basic verbosity  
ansible-playbook -i hosts.ini debug_validate.yml -v  
# Run with maximum verbosity  
ansible-playbook -i hosts.ini debug_validate.yml -vvv  
# Run in check mode  
ansible-playbook -i hosts.ini debug_validate.yml -check  
# Run debug tasks only  
ansible-playbook -i hosts.ini debug_validate.yml -tags "debug"
```

Program 6: Ansible Ecosystem and Reusable Content Management

Topics Included:

- **Ansible Galaxy** (Purpose, usage, finding roles)
- **Ansible Collections** (Purpose, benefits, basic usage)
- **Ansible Playbook Includes and Imports** (Differences, usage)

Rationale: Ansible Galaxy and Ansible Collections are both part of Ansible's ecosystem for managing reusable content (roles and collections). Galaxy is a hub for finding roles, while Collections extend this concept with modular, reusable code including roles, modules, and plugins. Playbook Includes and Imports tie into this by allowing you to integrate Galaxy roles or Collection content into your playbooks dynamically (`import_role`, `include_tasks`) or statically. s Program Example: A playbook that pulls a role from Galaxy (e.g., configuring an Nginx server), uses a Collection for additional utilities (e.g., `community.general`), and demonstrates `import_role` vs. `include_tasks` for modularity.

Program:

```
- name: Deploy Nginx using Galaxy Role and  
Collections  
  hosts: webserver  
  tasks:  
    - name: Import Nginx role from Galaxy  
      ansible.builtin.import_role:  
        name: geerlingguy.nginx # Fetched  
via ansible-galaxy  
    - name: Use Collection module for  
additional setup  
      community.general.package_facts:  
        manager: apt  
    - name: Include dynamic tasks  
      ansible.builtin.include_tasks:  
setup_firewall.yml
```

Program 7: Data Manipulation and Dynamic Playbooks

Topics Included:

- **Ansible Filters** (Basic usage, data manipulation)
- **Ansible Lookup Plugins** (Basic understanding, usage)
- **Ansible Playbook Variables Precedence** (Understanding the order)

Rationale: Filters and Lookup Plugins are tools for manipulating and retrieving data dynamically within playbooks. Filters transform data (e.g., `| json_query`), while Lookups fetch external data (e.g., `lookup('file', 'path')`). Variables Precedence is critical here because it determines how variables (used in filters or lookups) are overridden or prioritized (e.g., `playbook vars` vs. `role vars`). Program Example: A playbook that reads data from a file using a lookup, manipulates it with filters, and respects variable precedence for customization.

Program:

```
- name: Process server data dynamically  
  hosts: all  
  vars:  
    default_port: 80  
  tasks:  
    - name: Read config from file using lookup  
      ansible.builtin.set_fact:  
        config_data: "{{ lookup('file',  
'config.json') | from_json }}"  
    - name: Filter and transform data  
      ansible.builtin.debug:  
        msg: "Server: {{  
config_data.servers | map(attribute='name') |  
join(', ') }}"  
    - name: Show variable precedence (playbook  
vars override defaults)  
      ansible.builtin.debug:  
        msg: "Port: {{ port |  
default(default_port) }}"
```

Program 8: Robust Automation with Error Handling and Scaling

Topics Included:

- **Ansible Dynamic Inventory** (Basic concept, benefits)
- **Ansible Error Handling** (`ignore_errors`, `failed_when`)
- **Ansible Forks** (Basic understanding)

Rationale: Dynamic Inventory allows Ansible to adapt to changing environments (e.g., cloud instances), which pairs well with Forks for parallel execution across multiple hosts. Error Handling ensures robustness by managing failures (e.g., ignoring non-

critical errors or defining custom failure conditions).

Program Example: A playbook that uses a dynamic inventory (e.g., AWS EC2), handles errors gracefully, and scales with forks.

Program:

```
- name: Manage cloud servers with error handling
  hosts: all
  # Dynamic inventory assumed (e.g., ec2.py script)
  forks: 10 # Parallel execution
  tasks:
    - name: Install package with error handling
      ansible.builtin.package:
        name: httpd
        state: present
        ignore_errors: yes # Continue despite failures
    - name: Check service status
      ansible.builtin.command: systemctl status httpd
      register: result
      failed_when: "'running' not in result.stdout" # Custom failure condition
    - name: Debug result
      ansible.builtin.debug:
        msg: "Service is {{ 'up' if 'running' in result.stdout else 'down' }}"
```

Program 9: Controlled Deployment with Rolling Updates

Topics Included:

- **Ansible Rolling Updates (serial)**

Rationale: Rolling Updates (serial) is a standalone but critical concept for managing deployments in production environments, ensuring minimal downtime by updating hosts in batches. This can be a dedicated program as it's often used independently or combined with other features (e.g., error handling from Set 3).

Program Example: A playbook that updates a web application across multiple servers in batches.

Program:

```
- name: Perform rolling update on web servers
  hosts: webservers
  serial: 2 # Update 2 hosts at a time
  tasks:
    - name: Update application package
      ansible.builtin.package:
        name: myapp
        state: latest
    - name: Restart service
      ansible.builtin.service:
        name: myapp
        state: restarted
    - name: Verify application
      ansible.builtin.uri:
        url: "http://{{ inventory_hostname }}/health"
        status_code: 200
```